



Universitat
de les Illes Balears

Machine Learning

Lesson 2: Supervised Learning

Linear Models (LMS, Logistic Regression, Perceptron)

Remember: a simple example...

<i>Size (feet²)</i>	<i>Number of bedrooms</i>	<i>Number of floors</i>	<i>Age of home (years)</i>	<i>Price (\$1000)</i>
2104	5	1	45	460
1416	3	2	40	232
1534	3	2	30	315
852	2	1	36	178
...

... for dataset notation

$$\text{Dataset} = \{(\mathbf{x}^{(i)}, y^{(i)}); i = 1, \dots, m\}$$


$$(\mathbf{x}^{(i)}, y^{(i)}) = \text{Training example}$$

$$\mathbf{x}^{(i)} = \text{"input" variable (features)}, \mathbf{x}^{(i)} = (x_1^{(i)}, x_2^{(i)}, \dots, x_n^{(i)})$$


$$y^{(i)} = \text{"output" variable (target)}, y \in \mathcal{Y}$$

What price...?

Size (feet ²)	Price (\$) in 1000's
2104	460
1416	232
1534	315
852	178
...	...




ML Model: regression



Linear regression

Hypothesis: the model is linear

$$h_{\theta}(x) = \theta_0 + \sum_{j=1}^n \theta_j x_j$$



Case $n=1$:


$$h_{\theta}(x) = \theta_0 + \theta_1 x_1$$

Linear regression

Hypothesis: the model is linear

$$h_{\theta}(x) = \theta_0 + \sum_{j=1}^n \theta_j x_j$$

is called the **bias** because it represents the intercept or prior (prejudice)



Linear regression

Hypothesis: the model is linear

$$h_{\theta}(x) = \theta_0 + \sum_{j=1}^n \theta_j x_j$$

n = number of input dimensions

Generalization

(in matrix notation)

$$h_{\theta}(x) = \theta^T x$$

$x_0=1$ (Intercept term)

Loss function

Given a training set, how do we learn the parameters?

Basic idea: to make h_{θ} close to y


We've to define a function that measures, for each parameters' values, how close the $h_{\theta}(x^{(i)})$'s are to the corresponding $y^{(i)}$'s


Quadratic loss function

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

Case n=1:

$$h_{\theta}(x) = \theta_0 + \theta_1 x_1$$







X	Y	$h(\theta_0=0, \theta_1=0.5)$	$MSE = (h(\theta=0.5) - y)^2$
1	1		
20	12		

Quadratic loss function: example



$h_{\theta}(x) = \theta_1 x_1$
(for a fixed θ_1 , this is a function of x)



$J(\theta_1)$
(function of the parameter θ_1)




Quadratic loss function: example



Quadratic loss function: example


$$h_{\theta}(x) = \theta_1 x_1$$

(for a fixed θ_1 , this is a function of x)





$$J(\theta_1)$$

(function of the parameter θ_1)



Quadratic loss function: example



Analytic Linear regression

Dataset = X, \vec{y} , where X is a matrix $m \times (n + 1)$

$$J(\theta) = \frac{1}{2m} (X\theta - y)^T (X\theta - y)$$

$$\vec{y} = X\theta$$

$$X\theta = \vec{y}$$

$$X^T X \theta = X^T \vec{y}$$

$$\theta = (X^T X)^{-1} X^T \vec{y}$$


Size (feet ²)	Number of bedrooms	Number of floors	Age of home (years)	Price (\$1000)
2104	5	1	45	460
1416	3	2	40	232
1534	3	2	30	315
852	2	1	36	178
...

[full derivation](#)

LMS learning algorithm

Loss function: Quadratic

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$




Goal: To minimize $J(\theta)$


Gradient descent

$$\theta_j := \theta_j - \alpha \cdot \frac{\partial}{\partial \theta_j} J(\theta)$$



Learning rate





Gradient descent: basic idea





Gradient descent: Linear





Gradient descent: Linear



Gradient descent: Linear



Gradient descent: Linear



LMS algorithm: update rule

$$\theta_j := \theta_j - \alpha \cdot \frac{\partial}{\partial \theta_j} J(\theta) \longrightarrow \theta_j := \theta_j - \alpha \cdot (h_\theta(x) - y)x_j$$

Exercise: Obtain the gradient for the quadratic cost function of:

$$h_\theta(X) = \theta_0 + \theta_1 X$$

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)})^2$$

This rule (also known as Widrow-Hoff learning rule) has several properties that seem natural and intuitive. For instance, the magnitude of the update is proportional to the error term $(y - h_\theta(x))$; thus, for instance, if we are encountering a training example on which our prediction nearly matches the actual value, then, we find that there is little need to change the parameters; in contrast, a larger change to the parameters will be made if our prediction has a large error.

Exercise 2. Using the Gradient derived above, calculate the first iteration of the gradient descent


Initial Model

$$\hat{y} = \theta_0 + \theta_1 x$$

$$\begin{aligned}\theta_0 &= 2 \\ \theta_1 &= 1 \\ \alpha &= 0.1\end{aligned}$$

Grad Desc Algorithm

$$\theta_j := \theta_j - \alpha \cdot \frac{\partial}{\partial \theta_j} J(\theta)$$



Sample	X_0	X_1	y_real	θ_0^j	θ_1^j	y_pred	error_term	Grad_ θ_0	Grad_ θ_1	Learn_rate	θ_0^{j+1}	θ_1^{j+1}
0	1.0	2.50	5.58	2	1							
1	1.0	1.86	5.30									
2	1.0	2.65	6.96									
3	1.0	3.52	8.24									
4	1.0	1.77	5.38									

Learning rate


$$\theta_j := \theta_j - \alpha \cdot \frac{\partial}{\partial \theta_j} J(\theta)$$

Learning rate

In Practice:


If α is too small, gradient descent can be slow.

If α is too large, gradient descent can overshoot the minimum. It may fail to converge, or even diverge



Gradient descent can be susceptible to **local minima** in general, the optimization problem we have posed here for linear regression has only one global, and no other local, optima; thus gradient descent always converges (assuming the learning rate α is not too large) to the global minimum.

As we approach a local minimum, gradient descent will automatically take smaller steps, being $[0,1]$ smooths the derivative by reducing the jump stride.



Learning rate


$$\theta_j := \theta_j - \alpha \cdot \frac{\partial}{\partial \theta_j} J(\theta)$$

Learning rate

In Practice:


If α is too small, gradient descent can be slow.

If α is too large, gradient descent can overshoot the minimum. It may fail to converge, or even diverge



Gradient descent can be susceptible to **local minima** in general, the optimization problem we have posed here for linear regression has only one global, and no other local, optima; thus gradient descent always converges (assuming the learning rate α is not too large) to the global minimum.

As we approach a local minimum, gradient descent will automatically take smaller steps, being $[0,1]$ smooths the derivative by reducing the jump stride.



LMS algorithm: incremental rule

```
Loop{  
    for i=1 to m, {  
  
         $\theta_j := \theta_j + \alpha (y^{(i)} - h_{\theta}(x^{(i)})) x_j^{(i)}$   
  
    }  
}
```

When the training set is large, **stochastic gradient descent** is often preferred over batch gradient descent

Whereas batch gradient descent has to scan through the entire training set before taking a single step—a costly operation if m is large—stochastic gradient descent can start making progress right away, and continues to make progress with each example it looks at. Often, stochastic gradient descent gets θ “close” to the minimum much faster than batch gradient descent. (Note however that it may never “converge” to the minimum, and the parameters will keep oscillating around the global minimum; but in practice, most of the values near the minimum will be reasonably good approximations to the true minimum.

LMS algorithm: batch rule


epoch

Repeat until convergence {



$$\theta_j := \theta_j + \alpha \sum_{i=1}^m (y^{(i)} - h_{\theta}(x^{(i)})) x_j^{(i)}$$

}

(for every j)



Stochastic Gradient Descent with Batch size “1”



Gradient descent


minimize an objective function $J(\theta)$ by updating the parameters in the opposite direction of the gradient of the objective function.

The learning rate determines the size of the steps taken to reach the minimum

- Batch gradient descent: all training observations utilized in each iteration
- SGD: one observation per iteration
- Mini batch gradient descent: size of about 50 training observations for each iteration

LMS in practice

It is necessary to display the cost function





Simultaneous update

$$\begin{aligned}\text{temp0} &:= \theta_0 - \alpha \frac{\partial}{\partial \theta_0} J(\theta_0, \theta_1) \\ \text{temp1} &:= \theta_1 - \alpha \frac{\partial}{\partial \theta_1} J(\theta_0, \theta_1) \\ \theta_0 &:= \text{temp0} \\ \theta_1 &:= \text{temp1}\end{aligned}$$

No. of iterations

Data scaling /
normalization



Summary of Linear models

Linear regression


Hypothesis:

the model is linear

$$h_{\theta}(x) = \sum_{j=0}^n \theta_j x_j$$


Gradient descent

Idea: to make h_{θ} close to y by means minimizing a cost function



LMS

Using the gradient to find the minimum (batch & incremental)



Linear regression revisited

Hypothesis: the model is linear

n = number of input dimensions

$$h_{\theta}(x) = \sum_{j=0}^n \theta_j x_j$$

Generalization

(in matrix notation)

Case $n=1$:

$$h_{\theta}(x) = \theta^T x$$

$x_0=1$ (Intercept term)


$$h_{\theta}(x) = \theta_0 + \theta_1 x_1$$

Bias-Variance Tradeoff

A learning algorithm is **biased** for a particular input (x) if, when trained on different data sets, it is systematically incorrect when predicting the correct output for x .

Therefore, the **bias** is an error from erroneous assumptions in the learning algorithm.

High bias can cause an algorithm to miss the relevant relations between features and target outputs (**underfitting**).



$$\theta_0 + \theta_1 x$$


Bias
(Underfitting)

Bias-**Variance** Tradeoff

A learning algorithm has high **variance** for a particular input x if it predicts different output values when trained on different training sets.

The **variance** is an error from sensitivity to small fluctuations in the training set.

High variance can cause an algorithm to model the random noise in the training data, rather than the intended outputs (**overfitting**).



$$\theta_0 + \theta_1 x + \theta_2 x^2 + \theta_3 x^3 + \theta_4 x^4$$


Variance
(Overfitting)

Bias-Variance Tradeoff

A learning algorithm is **biased** for a particular input (x) if, when trained on different data sets, it is systematically incorrect when predicting the correct output for x .

Therefore, the **bias** is an error from erroneous assumptions in the learning algorithm.


High bias can cause an algorithm to miss the relevant relations between features and target outputs (**underfitting**).



$$\theta_0 + \theta_1 x$$

Bias
(Underfitting)

Bias-Variance Tradeoff



Bias-Variance Tradeoff

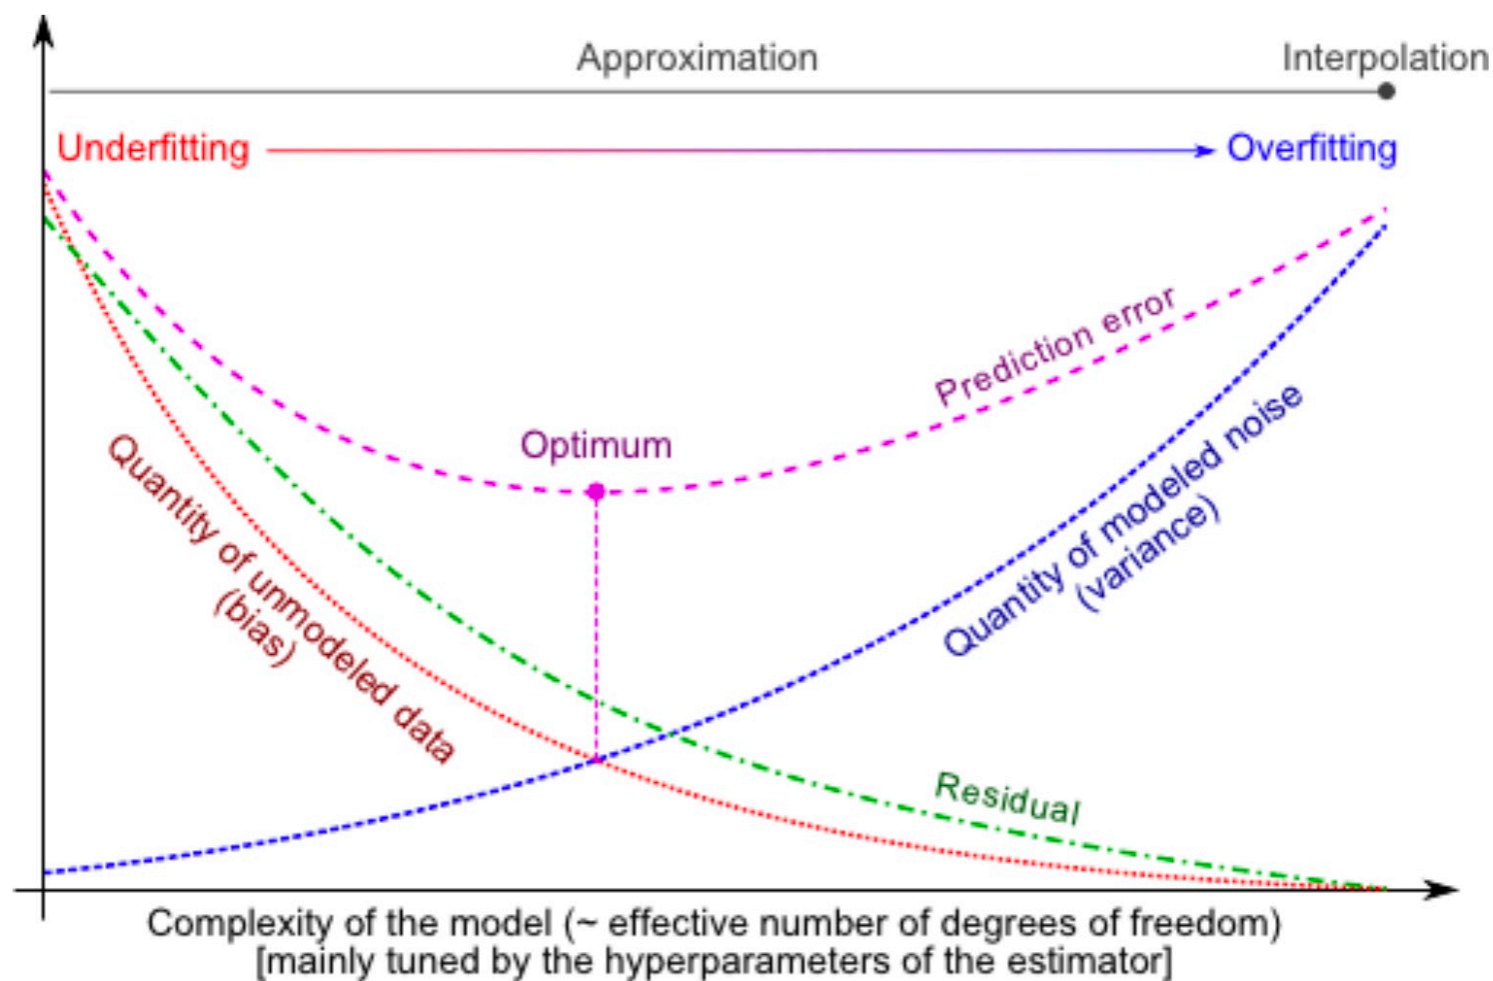
$$\begin{aligned} Err(x) = \mathbb{E} \left[(y_{test} - \hat{y})^2 \right] &= \left(E[\hat{f}(x)] - f(x) \right)^2 + E \left[\left(\hat{f}(x) - E[\hat{f}(x)] \right)^2 \right] \\ &+ \sigma_e^2 \end{aligned}$$

Bias: La diferencia entre el valor predicho por el modelo y el valor real.

Varianza: El error producido debido a la sensibilidad del modelo con respecto a los datos de entrenamiento.

Error irreducible: ruido que no conseguiremos modelizar

Bias-Variance Tradeoff




Bias-Variance Tradeoff

Alto bias: Simplificaciones sobre la forma de los datos, *menos flexibles*: Linear models

Baja Varianza: Pequeños cambios en las predicciones: Linear models

Bajo bias: Pocos supuestos sobre la forma de los datos, *más flexibles*: Decision trees, knn, ...

Alta Varianza: Grandes cambios en las predicciones: Decision trees, knn, ...



Bias-Variance Tradeoff

La única forma de comprobar que nuestro modelo generaliza bien es guardando una parte de los datos para medir el error *out-of-sample*, es decir, el error sobre los datos que no ha visto el modelo durante la fase de entreno.

Con tal de poder elegir el mejor modelo y sus parametros se dividirá inicialmente el dataset en dos subsets:

- subset de entrenamiento: normalmente se utiliza el 80% del tamaño de la muestra
- subset de testeo: el 20% restante.

